Arguing for Decisions: A Qualitative Model of Decision Making

Blai Bonet

Dpto. de Computación Universidad Simón Bolívar Aptdo. 89000, Caracas 1080-A Venezuela

Abstract

We develop a qualitative model of decision making with two aims: to describe how people make simple decisions and to enable computer programs to do the same. Current approaches based on Planning or Decision Theory either ignore uncertainty and tradeoffs, or provide languages and algorithms that are too complex for this task. The proposed model provides a language based on rules, a semantics based on high probabilities and lexicographical preferences, and a transparent decision procedure where reasons for and against decisions interact. The model is no substitute for Decision Theory, yet for decisions that people find easy to explain it may provide an appealing alternative.

1 INTRODUCTION

In this paper we develop a qualitative model of decision making with two aims: to describe how people make simple, everyday decisions and to enable computer programs to do the same. Current approaches based on Planning [Weld, 1994] or Decision Theory [Raiffa, 1970] either ignore uncertainty and tradeoffs, or provide languages and algorithms that are too complex for this task. The model proposed provides a simple language based on rules, a semantics based on high probabilities and lexicographical preferences, and a transparent decision procedure where reasons for and against decisions interact.

The model is closely related to other qualitative abstractions of Decision Theory (e.g., [Pearl, 1993; Boutilier, 1994; Wilson, 1995]), yet it introduces as-

Mailing address from US and Europe: Hector Geffner, Bamco CCS 144-00, P.O.BOX 02-5322, Miami Florida 33102-5322, USA. E-mail: {bonet,hector}@usb.ve.

Hector Geffner*

Dpto. de Computación Universidad Simón Bolívar Aptdo. 89000, Caracas 1080-A Venezuela

sumptions that aim to account for the way decisions are made in simple settings. In the proposed model, like in the findings of [Shafir et al., 1993; Hogarth and Kunreuther, 1995], the reasons for decisions play a central role. The result is an efficient, 'anytime' decision procedure, which is easy to justify and explain.

The paper is organized as follows. First we introduce the representation language (Section 2), the decision procedure (Section 3) and the semantics (Section 4). Then we discuss the relation to Decision Theory (Section 5), sensitivity issues (Section 6), extensions (Section 7), and related work (Section 8).

2 LANGUAGE

Models in the proposed framework contain four parts (see Fig. 1).

- 1) a set of input propositions and observations defining the possible input situations,
- 2) a set of goals and goals priorities defining the output situations,
- 3) a set of actions and action rules defining how input situations are mapped to output situations, and
- 4) plausibility measures defining the plausibility of the input situations

For example, a situation in which one has to decide whether to study for an exam or go to the beach can be modeled as:

study ∧ get-it ⇒ pass-exam go-beach ∧ ¬rain ⇒ enjoy-beach unlikely rain ; plausible get-it

Here study and go-beach are the possible actions, rain and get-it are the input propositions and pass-exam and enjoy-beach are the positive goals in that order of importance.

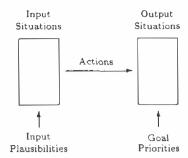


Figure 1: The Decision Situation

2.1 INPUT SITUATIONS

The input situations or states stand for the possible truth assignments to the input propositions and represent the context of the decision. With each input proposition y we associate a boolean variable Y such that y stands for Y = true and $\neg y$ stands for Y = false. We also use the notation Y to denote any of the literals y or $\neg y$ associated with the variable Y, and $\sim Y$ to denote its complement. The observations are input literals that have been found to be true.

2.2 GOALS

The goals stand for states of affairs that we care about. The positive goals are the ones that we want to achieve: getting a good job, enjoying a good day at the beach, watching a good movie, etc. The negative goals are the ones that we want to avoid: being dead, being thirsty, hurting people, missing an appointment, etc.

In this model, the goals are represented by literals (different from the input literals; yet see Section 7) denoted by symbols like $\mathbf{x}, \mathbf{x}', \ldots$. The set of all goals is denoted by the letter G while the set of positive and negative goals by G^+ and G^- respectively. We use the words goal literals to refer to goals or their complements.

As the example above suggests, some goals are more important than others: getting a good job is more important than watching a good movie, being not dead is more important than missing an appointment, etc. We represent the relative importance of goals by integers: the higher the integer, the higher the importance of the goal. We call such integers the priority of the goals, and write $\mathbf{x} \in G_i$ to say that the priority of goal \mathbf{x} is i (we will also write $\mathbf{x} \in G_i^+$ or $\mathbf{x} \in G_i^-$ when we want to make explicit the polarity of the goal as well). Schank and Abelson [1977] and Slade [1992] provide an interesting analysis of different types of goals (e.g., satisfaction, enjoyment, achievement, preservation, etc.) and their relative priorities.

Goal priorities are related to goal utilities, yet as de-

grees of *importance*, we assume that goal priorities combine as follows:

Assumption 1 Higher priority positive (negative) goals should be pursued (avoided)—even at expense of lower priority goals—except when success is deemed unlikely.

This is a fundamental assumption in the model and says to focus on the actions that serve the most important goals, ignoring unlikely possibilities. This is not always a reasonable thing to do (see Section 6) but seems appropriate in the context of simple, everyday decisions. As we will see, this assumption will allow us cast the decision process as an argumentation process where reasons for and against decisions interact.

Because goals are important, we also assume that goals that are not said to be true explicitly,¹, are not true:²

Convention 1 Goals are assumed not true by default.

Due to this convention, there is a difference in this model between declaring x as a positive goal and its complement $\sim x$ as a negative goal. Even though in both cases we will try to achieve x and avoid $\sim x$, in the first case $\sim x$ will be assumed true by default while in the second x will. In line with this convention we require that if x is a goal, $\sim x$ is not.

2.3 ACTIONS AND ACTION RULES

The third component of the model are the actions and the action rules. Action rules map input situations (truth assignments to the input literals) to output situations (truth assignments to the goal literals). They are expressed by means of expressions of the form: $A \wedge C \Rightarrow \mathbf{x}$, where A is an action symbol, C is a conjunction of input literals, and \mathbf{x} is a goal literal (action symbols are distinct from input and output symbols).

Action rules are default rules in the sense that \mathbf{x} is normally true after doing A when C is true. Each action rule has a priority or strength measure represented by a non-negative integer; the higher the number, the higher the priority. These priorities will be used to disambiguate conflicts among rules; e.g., to make a rule $A \wedge C \Rightarrow \mathbf{x}$ override a conflicting rule $A \wedge C' \Rightarrow \sim \mathbf{x}$ of lower priority. Unless otherwise specified, all rules are assumed to have priority zero (lowest priority).

Action rules which do not involve any actions, like knows-a-lot \Rightarrow pass-exam, will be interpreted as

¹Actually, there is no way to explicitly say that a goal is true in this language, yet see Section 7 for extensions that do.

²We distinguish 'assumptions' from 'conventions' to emphasize that the latter are just a matter of convenience; they are not built into the model like the former.

abbreviations of rules involving the special action do-nothing, e.g., do-nothing \land knows-a-lot \Rightarrow pass-exam. The action do-nothing is assumed always present and represents the choice of not taking any (other) action. As for other actions, we can also have rules involving the symbol do-nothing explicitly (e.g., if a person is seriously injured and you do nothing, the person may die, etc.).

2.4 INPUT PLAUSIBILITIES

The last component of the model are the plausibilities of the input propositions. For that we allow the following type of statements for any input literal Y: 'likely Y', 'plausible Y' and 'unlikely Y'. Intuitively, these statements rank the prior probability of Y in decreasing order, with the first and last denoting probabilities that are very close to 1 and 0 respectively. The meaning of these statements will be made precise in terms of Spohn's [1988] κ -functions. For the user's convenience we assume that:

Convention 2 When the input statements do not contain information about the plausibility of an input literal Y, Y is assumed plausible.

2.5 EXAMPLE

The situation of going for the newspaper with or without the umbrella can be modeled in this language by means of action rules like:

We also have to say that the possible actions are go-with-umbrella and go-without-umbrella and that the goals (and their polarities and priorities) are carry $\in G_1^-$, wet $\in G_2^-$ and newspaper $\in G_3^+$ (i.e., getting the newspaper is the most important goal, and avoiding getting wet is more important than avoiding having to carry an umbrella).

From the conventions above, it is implicit that rain is plausible (Convention 2), that each of the goals newspaper, wet and carry are true only when a rule asserting the goal is applicable (Convention 1), and that the action do-nothing does not achieve any goal.

3 REASONS FOR DECISIONS

We present now a *mechanism* for deciding which action to choose in a given context. The mechanism is based on the interplay of reasons. The procedure is efficient

and easy to justify and explain. We start defining the reasons for decisions.

Basically, we will say that a positive (negative) goal x provides a reason for (against) action A when the action A contributes to the truth of x. The polarity of this reason is the polarity of the goal (positive or negative); the importance of the reason is the priority of the goal $(0, 1, \ldots, N)$; and the strength of the reason is the measure to which the action contributes to the truth of the goal.

More formally, let us say that a literal \mathbf{Y} is likely when \mathbf{Y} is an observation or when the information provided by the user contains likely \mathbf{Y} or unlikely $\sim \mathbf{Y}$, and that \mathbf{Y} is unlikely when its complement is likely, and plausible when \mathbf{Y} is neither likely nor unlikely. Similarly, let us say that a rule $A \wedge C \Rightarrow \mathbf{X}$ is likely when each conjunct in C is likely, that is unlikely when some conjunct in C is unlikely, and that is plausible when it is neither likely nor unlikely. Then, the reasons for decisions and their strengths are defined as follows:

Definition 1 A positive (negative) goal x provides a strong reason for (against) an action A when some rule $A \wedge C \Rightarrow x$ is likely and no rule of the form $A \wedge C' \Rightarrow x$ with equal or higher priority is either likely or plausible.

Definition 2 A positive (negative) goal x provides a weak reason for (against) an action A when it does not provide a likely reason for A and yet some rule $A \wedge C \Rightarrow x$ is either likely or plausible, and no higher priority rule $A \wedge C' \Rightarrow \sim x$ is likely.

Definition 3 A positive (negative) goal x provides a empty reason for (against) an action A when it does not provide a strong or weak reason for (against) A.

As an illustration, the goal newspaper provides a strong reason for go-with-umbrella and for go-without-umbrella; wet provides a weak reason against go-without-umbrella, and carry provides a weak reason against go-with-umbrella. Likewise, each of these goals provide empty reasons for or against do-nothing.

Clearly, decisions over a single goal can be taken by considering the strength and polarity of the reasons involved.

Definition 4 An action A is better than an action B over a positive goal x when x provides a stronger reason for A than for B. Likewise, A is better than B over a negative goal x if x provides a stronger reason against B than against A.

³Strong reasons are stronger than weak reasons, and weak reasons are stronger than empty reasons.

When there are many goals involved, the more important goals are considered first:

Definition 5 An action A is better than an action B, written A > B, when A is better than B over a goal x and B is no better than A over any goal x' as important as or more important than x.

The overall best actions are the actions that are no worse than any other action. We can test whether an action A is better than an action B by invoking the procedure better?(A, B, i) that iteratively checks whether A gets more compelling reasons than B from goals in G_i , where i is a priority level initially set to the top priority N. Indeed, better? must return no when some positive (negative) goal provides a stronger reason for B (A) than for A (B); yes when the opposite is true, and must call itself with the value of i decreased when neither condition holds, returning no when i < 0. In the worst case, the complexity of this method is:

Proposition 1 The best actions can be computed in this way in time proportional to $A^2 \times R$, where A is the number of actions and R is the total number of rules.

This complexity of this method is moderate, yet a more efficient procedure can be used when goals are linearly ordered (i.e., when no pair of goals have the same priority). If \mathcal{A} stands for the set of all actions and i is a priority level (initially set to N), $select(\mathcal{A}, i)$ can compute the best actions by retaining in \mathcal{A} , in each iteration, only the actions that get the strongest (weakest) reason from the single positive (negative) goal in G_i . This iteration terminates when i < 0 or when \mathcal{A} becomes a singleton.

In the example above, do-nothing is pruned from \mathcal{A} in the first iteration because it only gets an empty reason from the positive goal newspaper. In the second iteration, go-without-umbrella is also pruned as it gets a strong negative reason from the goal wet. The action go-with-umbrella then is the single best action as it is the only action left in \mathcal{A} .

3.1 EXAMPLE

Consider whether to approach some animal, e.g., a dog, that we don't know whether it is aggressive or not:

$$\begin{array}{ccc} & \text{approach} & \Rightarrow & \text{satisfy-curiosity} \\ & \text{approach} \land \text{aggressive} & \Rightarrow & \text{get-hurt} \end{array}$$

In this case, satisfy-curiosity is a low priority positive goal and get-hurt is a high priority negative goal. Given no other information, aggressive

is assumed plausible by convention, and hence, the action approach gets a strong positive reason from satisfy-curiosity and a weak negative reason from get-hurt. Yet since get-hurt is the most important goal, the action approach is rejected for do-nothing which gets no (weak or strong) negative reasons. Note, however, that if observations lead to us revise the chances of aggressive to unlikely, the preferences would get reversed.

4 SEMANTICS

The semantics will make precise the meaning of all the constructs in the model and will provide an independent criterion for assessing the decision procedures above. In Decision Theory, actions A are ranked by their expected utility:

$$EU(A) = \sum_{s} P(s) \sum_{s'} P_A(s'|s) U(s')$$
 (1)

where the s and s' denote the input and output states respectively. Here we use an approximation of this criterion with Spohn's [1988] κ -functions in place of probabilities, and lexicographical orderings in place of utility functions.

4.1 BELIEFS

Spohn [1988] describes a model for uncertain reasoning that combines the main intuitions underlying probability theory (context dependence, conditionalization, etc.) with the notion of plain beliefs (see also [Goldszmidt and Pearl, 1992]). Beliefs in Spohn's model are represented by means of a function κ that assigns a non-negative integer measure to each world w and that satisfies the following calculus:⁴

$$\kappa(p) \in [0, \infty], \ \kappa(p) = \min_{w \models p} \kappa(w)$$
(2)

$$\kappa(p \vee \neg p) = 0$$
, $\kappa(p|q) = \kappa(p \wedge q) - \kappa(q)$ (3)

This calculus is structurally similar to the calculus of probabilities with products replaced by sums and sums replaced by minimizations. Spohn indeed showed that the first can be understood as an abstraction of the second with κ measures standing for order-of-magnitude probabilities.

Lower κ measures stand for higher probabilities and higher κ measures stand for lower probabilities. Spohn indeed refers to the κ measures as degrees of surprise or disbelief, hence regarding a proposition p as plausible or believable when $\kappa(p) = 0$ and as unlikely or disbelieved when $\kappa(p) > 0$. In particular, since the axioms rule out two complementary propositions from

 $^{^4\}kappa(p)=\infty$ when p is unsatisfiable.

being disbelieved at the same time, a proposition p is accepted or believed when its negation is disbelieved (i.e., when $\kappa(\neg p) > 0$).

4.1.1 INPUT BELIEFS

We use Spohn's κ functions to formalize the beliefs $\kappa(\mathbf{Y})$ in the inputs propositions and the beliefs $\kappa_A(\mathbf{x}|s)$ in the output propositions given an input state and action. For the inputs beliefs, we assume that:⁵

Assumption 2 Input variables are assumed to be independent.

This means that, in analogy to probabilities, the belief in an input state is the aggregation of the beliefs in the input literals true in that state:

$$\kappa(\mathbf{Y}_1, \dots, \mathbf{Y}_n) = \sum_{i=1}^n \kappa(\mathbf{Y}_i) \tag{4}$$

These beliefs in turn are provided by the user or assumed by default (Convention 2):

$$\kappa(\mathbf{Y}_i) = \begin{cases} 1 & \text{if unlikely } \mathbf{Y}_i \text{ or likely } \sim \mathbf{Y}_i \\ 0 & \text{otherwise} \end{cases}$$
 (5)

These two equations determine the prior plausibility $\kappa(s)$ of any input state s completely. The posterior plausibility $\kappa(s|\mathsf{obs})$ given a set of observations (input literals) can be derived from (3) as $\kappa(s|\mathsf{obs}) = \kappa(s) - \kappa(\mathsf{obs})$ if s satisfies obs , and ∞ otherwise.

4.1.2 GOAL BELIEFS

For any goal literal x, the plausibility of x given an action A and an input state s is expressed by the equation:

$$\kappa_A(\mathbf{x}) = \min_{s} \left(\kappa(s) + \kappa_A(\mathbf{x}|s) \right) \tag{6}$$

which is the qualitative version of the equation $P_A(\mathbf{x}) = \sum_s P(s) P_A(\mathbf{x}|s)$.

From Equations 4 and 5, we know how to determine the plausibilities $\kappa(s)$; we are thus left to determine the conditional plausibilities $\kappa_A(\mathbf{x}|s)$. These plausibilities will be extracted from the rules that are applicable in the state s that relate A to \mathbf{x} .

Let r(A) denote the set of action rules involving the action A. Then we say that the input state s supports a literal x when there is a rule in r(A) whose consequent is x, whose conditions are true in s, and for which all conflicting rules in r(A) whose conditions are true in s have equal or lower priority.

The plausibilities $\kappa_A(\mathbf{x}|s)$ are then defined to capture Convention 1 (goals are assumed false by default) and

the intuition that supported literals should be either likely or plausible (they have a justification):

$$\kappa_{A}(\mathbf{x}|s) = \begin{cases}
0 & \text{when } \mathbf{x} \text{ is supported by } s, \text{ or when} \\
-\mathbf{x} \text{ is a goal not supported by } s \\
1 & \text{otherwise}
\end{cases}$$
(7)

Equations 4-7 determine the measures $\kappa_A(\mathbf{x})$ for any value \mathbf{x} and any action A. When input observations obs are gathered, the conditional measure $\kappa_A(\mathbf{x}|\text{obs})$ can be obtained by replacing the prior plausibility $\kappa(s)$ in Equation 6 by the posterior plausibility $\kappa(s|\text{obs})$.

4.1.3 EXAMPLE

We illustrate these definitions in the newspaper example. Because the only input variable is rain, which is assumed plausible by default, the input states are $s = \{\text{rain}\}\ \text{and}\ s' = \{\text{-rain}\}\ \text{with}\ \kappa(s) = \kappa(s') = 0.$

If A, B and C denote the actions go-without-umbrella, go-with-umbrella and do-nothing, the literals supported by each action in each input situation are:

	$s = \mathtt{rain}$	$s' = \neg rain$
A	$\{news, wet, \neg carry\}$	{news,¬wet,¬carry}
B	{news,¬wet,carry}	{news,¬wet,carry}
\overline{C}	$\{\neg news, \neg wet, \neg carry\}$	{¬news,¬wet,¬carry}

From this table and Equations 4-7, the plausibilities of all goal literals can be computed; e.g., $\kappa_A(\text{wet}) = \min\{\kappa(s) + \kappa_A(\text{wet}|s), \kappa(s') + \kappa_A(\text{wet}|s')\} = \min\{0 + 0, 0 + 2\} = 0.$

4.2 PREFERENCES OVER ACTIONS

To rank the actions, we define the qualitative utility of a goal x, written u(x), as:

$$u(\mathbf{x}) = polarity(\mathbf{x}) \times priority(\mathbf{x})$$
 (8)

Namely, for wet $\in G_2^-$, u(wet) = -2, while for newspaper $\in G_3^+$, u(newspaper) = 3.

Provided with these measures, we could define the qualitative expected utility of actions relative to a goal \mathbf{x} , following the methods in [Pearl, 1993] or [Wilson, 1995], e.g., setting it to $\max(0, u(\mathbf{x}) - \kappa_A(\mathbf{x}|\mathsf{obs}))$ when \mathbf{x} is positive. The problem with these schemes is that they impose a very strong requirement on the way utility measures are encoded so they can be added up, in the same scale, with κ measures (see [Wilson, 1995]).

Here we take a different approach which does not require goal priorities and plausibility judgements to be so calibrated. In the proposed scheme, only two things matters: the *ordinal* ranking of goals, and whether

⁵We will show how to relax this assumption in Section 7.

goals are deemed likely, unlikely or plausible. This is done by defining the *qualitative belief* in a goal literal **x** as:

$$b_A(\mathbf{x}) = \kappa_A(\neg \mathbf{x} | \mathsf{obs}) - \kappa_A(\mathbf{x} | \mathsf{obs}) \tag{9}$$

and defining the qualitative rank of an action A relative to a goal \mathbf{x} as:⁶

$$Q_{\mathbf{X}}(A) = sign(u(\mathbf{X})) \times sign(b_A(\mathbf{X})) \tag{10}$$

In other words, an action has a positive rank relative to a goal \mathbf{x} ($Q_{\mathbf{X}}(A)=1$) when it's likely to make the positive (negative) goal \mathbf{x} true (false); it has a negative rank ($Q_{\mathbf{X}}(A)=-1$) when it's likely to make the positive (negative) goal \mathbf{x} false (true); and it has an null rank otherwise ($Q_{\mathbf{X}}(A)=0$). Clearly,

Definition 6 An action A is preferred to an action B over a goal x if $Q_X(A) > Q_X(B)$.

In the presence of multiple goals, this ordering is extended by considering more important goals first:

Definition 7 An action A is preferred to an action B, written A > B, if A is preferred to B over some goal x, and B is not preferred to A over any goal x' with equal to or higher priority than x.

The overall optimal actions determined by this preference relation are closely related to the actions that result from the decision procedures based on rules (Section 3). Indeed, if we say that a theory is *positive* when the rules do not involve negative literals in their bodies we get that:

Proposition 2 The decision procedures based on interaction of reasons are sound and complete for positive theories.

The condition of positivity is required because the procedures do not reason by cases and thus cannot properly handle pairs of rules like $A \wedge p \Rightarrow \mathbf{x}$ and $A \wedge \neg p \Rightarrow \mathbf{x}$. Under this condition, Proposition 1 guarantees that the problem of identifying the best actions can be computed efficiently.

4.2.1 EXAMPLE

The table below summarizes the qualitative rank of the actions relative to each of the goals newspaper, wet and carry:

Action
$$\langle n, w, c \rangle$$
go-without-umbrella: $\langle +1, +0, +1 \rangle$
go-with-umbrella: $\langle +1, +1, -1 \rangle$
do-nothing: $\langle -1, +1, +1 \rangle$

The preferences among the actions are easy to visualize as they correspond to the lexicographical preferences among their corresponding vectors (in this case, no pair of goals have the same priority). The table makes evident that go-with-umbrella is the best action in this case. On the other hand, if rain were unlikely, the first entry below wet would become +1, and the best action would become go-without-umbrella.

5 RELATION TO DECISION THEORY

The optimal action A from a decision theoretic point of view is the action that maximizes the expected utility (Equation 1). The model can be understood as assuming that the utility function U(s) is additively decomposable as $U(s) = \sum_{\mathbf{X} \in s} U(\mathbf{X})$, where \mathbf{X} is the value of variable X in the output situation s, and that the utility of positive (negative) goals \mathbf{X} is a fixed positive (negative) value $U_{\mathbf{X}}$ and the utility for the negation of a goal is 0. From these assumptions, it is possible to show that Equation 1 can be expressed as:

$$EU(A) = \sum_{\mathbf{X} \in G} U_{\mathbf{X}} P_{A}(\mathbf{X})$$
 (11)

Furthermore, the model assumes that terms $U_{\mathbf{X}}P_A(\mathbf{x})$ make terms $U_{\mathbf{X}'}P_A(\mathbf{x}')$ negligible when $P_A(\mathbf{x}) \gg P_A(\mathbf{x}')$ ('unlikely scenarios are ignored') and that $|U_{\mathbf{X}}| \gg |U_{\mathbf{X}'}|$ when the priority of \mathbf{x} is higher than the priority of \mathbf{x}' ('low priority goals are traded by higher priority goals').

6 SENSITIVITY ISSUES

It is not hard to think of cases where the assumptions embedded in this model are not reasonable. Consider for example a situation in which a patient has a very serious disease which if not treated will result in his death. Moreover, there is only one possible treatment and such treatment does not always work, and in all cases it has undesirable side-effects like loosing hair, vomiting, etc.

$$\begin{array}{ccc} & \text{do-nothing} & \Rightarrow & \text{death} \\ & \text{treatment} \land \neg \text{effective} & \Rightarrow & \text{death} \\ & & \text{treatment} & \Rightarrow & \text{side-effects} \end{array}$$

Here the goals death and side-effects are both negative and the first is significantly more important than the second.

The atom effective provides the condition under which the treatment works. If the prior plausibility measure of effective is 0 (the treatment can plausibly work) the model recommends treatment. Yet if

⁶The function sign maps positive numbers into 1, negative numbers into -1, and 0 into 0.

the prior plausibility measure of effective is 1 (the treatment most likely will not work) the model will recommend to do nothing (i.e., the action do-nothing will be preferred to treatment).

One way to look at the second scenario is that the model prefers the lottery 'certain death' to the lottery 'certain side-effects and very likely death'. This preference, which is not reasonable, results from regarding unlikely scenarios as impossible ones. This assumption, in cases where important goals are at stake, is actually far from appropriate.

We can measure though how robust an optimal decision is by considering how it is affected by changes in the input parameters (goal priorities and input plausibilities).

Let us say that a goal x justifies the preference of action A over action B if A is preferred to B over x and yet A and B are equally preferred over all goals with higher priority than x.

For example, when the treatment is unlikely to work, the goal that justifies the decision do-nothing over treatment is side-effects. On the other hand, when the treatment can plausibly work, do-nothing becomes inferior to treatment because of the goal death. Since death is considerably more important than side-effects the proper selection of the parameter κ (effective) in this case is critical. More generally, when minimal changes in an input parameter lead to abrupt changes in the importance of the goals that are obtained the optimality of the decisions need to be reconsidered. This critical tradeoff can be detected in this model, yet the same model is not sufficiently expressive to resolve them. Often, however, there may be no reasonable ways for resolving such tradeoffs.

7 EXTENSIONS

The expressive power of the model is limited yet there are a number of extensions that can be accommodated.

First, we can relax the assumptions that input variables be independent by accommodating input rules in addition to action rules. These input rules will impose a causal structure on the input variables which can be interpreted as in [Goldszmidt and Pearl, 1992] or [Geffner, 1996a]. Semantically the only difference is in the determination of the plausibilities of the input state $\kappa(s)$.

Second, we can interpret the input and output situations as referring to the state of the world before and after the action. The values of variables that occur in both the inputs and the outputs can then be assumed to persist by default [Gelfond and Lifschitz, 1993;

Geffner, 1996b]. This can enable us to express sequential decision problems, where the choice of optimal actions is replaced by the choice of optimal action sequences.

In many cases, we may also need a way for representing and aggregating preferences among equally important goals. That is, two goals may be equally important and yet one may be preferred to the other; e.g., going to see the 'Knicks' vs. going to see the 'Mets'. A possible approach in this case is to express these preferences by means of integers and to aggregate such preferences by some form of weighted addition according to whether the goals are rendered likely, plausible or unlikely by the actions.

8 RELATED WORK

The proposed model for decisions is related to other qualitative abstractions of decision theory and to informal models of decisions based on the interplay of reasons.

Qualitative models of decision making have received considerable attention in recent years [Pearl, 1993; Boutilier, 1994; Dubois and Prade, 1995; Wilson, 1995]. All of these proposals have in common the use of qualitative measures for representing preferences and beliefs, yet compared to this work, few have placed emphasis on modeling (yet, see [Brewka and Gordon, 1995]) and in the mechanisms for computing and explaining decisions.

The work differs from [Pearl, 1993] and [Wilson, 1995] in the way utility ranks and κ measures are combined. Pearl and Wilson assume that these measures are calibrated so that they can be added up in the same scale. Thus, a likely world with utility rank 1 is deemed as good as an unlikely world with utility measure 2. Our choice here is different: our priority measures are completely ordinal and represent the importance of goals. Our criterion is that most important goals dominate less important goals except when the former are unlikely to be realized.

The two criteria can be usefully contrasted in the simple case in which there is a single positive goal \mathbf{x} involved. This scenario can be expressed in Pearl's and Wilson's framework by partioning the set of worlds into two sets: the worlds w^+ that satisfy \mathbf{x} , which get a utility rank $\mu(w^+)=1$, and the worlds w^- that do not satisfy \mathbf{x} , which get a utility rank $\mu(w^-)=0$. A weakness of Pearl's and Wilson's scheme is that they fail to prefer actions A that make \mathbf{x} likely $(\kappa_A(\neg \mathbf{x})>0)$ to actions B that make \mathbf{x} just plausible $(\kappa_B(\mathbf{x})=\kappa_B(\neg \mathbf{x})=0)$. Both actions get actually the same expected utility rank in their scheme. Inter-

estingly this is not solved when the worlds w^- that do not satisfy the goal are assigned a negative utility rank $\mu(w^-) = -1$. In that case, Pearl's and Wilson's schemes will label the actions B that make the goal x plausible, ambiguous. We, on the other hand, rank such actions below the actions A that make x likely, and above the actions C that make x unlikely.

The procedures considered in Section 3 are related also to informal models of decision based on the interplay of reasons. For example, when one action A gets either positive or negative (non-empty) reasons such that no other action gets (non-empty) reasons of the same importance, the action A can immediately be accepted, if the reasons are positive, and rejected, if the reasons are negative. These type of situations, where there are clear and compelling reasons for accepting or rejecting decisions, seem to be the ones people feel most comfortable with and have been studied in [Shafir et al., 1993].

Acknowledgments

We want to thank the anonymous UAI reviewers for useful comments.

References

- [Boutilier, 1994] C. Boutilier. Toward a logic for qualitative decision theory. In *Proceedings KR-94*, pages 75–86, 1994.
- [Brewka and Gordon, 1995] G. Brewka and T. Gordon. How to buy a porsche: An approach to defeasible decision making (extended abstract). In Workshops in Computational Dialectics, 1995.
- [Dubois and Prade, 1995] D. Dubois and H. Prade. Possibility theory as a basis for qualitative decision theory. In *Proceedings of IJCAI-95*, pages 1924–1930, 1995.
- [Geffner, 1996a] H. Geffner. A formal framework for causal modeling and argumentation. In *Proceedings* of Intl. Conf. on Formal and Applied Practical Reasoning. Springer-Verlag, 1996.
- [Geffner, 1996b] H. Geffner. A qualitative model for temporal reasoning with incomplete information. In Proceedings AAAI-96, Portland, Oregon, 1996. MIT Press.
- [Gelfond and Lifschitz, 1993] M. Gelfond and V. Lifschitz. Representing action and change by logic programs. J. of Logic Programming, 17:301-322, 1993.
- [Goldszmidt and Pearl, 1992] M. Goldszmidt and J. Pearl. Rank-based systems: A simple approach

- to belief revision, belief update and reasoning about evidence and actions. In *Proceedings KR-92*, pages 661-672, 1992.
- [Hogarth and Kunreuther, 1995] R. Hogarth and H. Kunreuther. Decision making under ignorance: Arguing yourself. J. of Risk and Uncertainty, 10:15-36, 1995.
- [Pearl, 1993] J. Pearl. From conditional oughts to qualitative decision theory. In *Proceedings UAI-93*, 1993.
- [Raiffa, 1970] H. Raiffa. Decision Analysis: Introductory Lectures on Choices under Uncertainty. Addison-Wesley, 1970.
- [Schank and Abelson, 1977] R. Schank and R. Abelson. Scripts, Plans, Goals and Understanding. Laurence Erlbaum Associates, Hillsdale, N.J., 1977.
- [Shafir et al., 1993] E. Shafir, I. Simonson, and A. Tversky. Reason-based choice. Cognition, 11:11-36, 1993.
- [Slade, 1992] S. Slade. An Interpersonal Model of Goal-Based Decision Making. PhD thesis, Department of Computer Science, Yale University, 1992.
- [Spohn, 1988] W. Spohn. A general non-probabilistic theory of inductive reasoning. In Proceedings 4th Workshop on Uncertainty, pages 315-322, St. Paul, 1988.
- [Weld, 1994] D. Weld. An introduction to least commitment planning. AI Magazine, 1994.
- [Wilson, 1995] N. Wilson. An order of magnitude calculus. In *Proceedings UAI-95*. Morgan Kaufmann, 1995.